

In [1]:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

In [3]:

```
!mkdir -p ~/.kaggle
!mv kaggle.json ~/.kaggle/
```

In [5]:

```
!kaggle datasets download -d salader/dogs-vs-cats
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /root/.kaggle/kaggle.json'

Downloading dogs-vs-cats.zip to /content

99% 1.05G/1.06G [00:11<00:00, 154MB/s]

100% 1.06G/1.06G [00:11<00:00, 98.3MB/s]

In [6]:

```
import zipfile
zip_ref = zipfile.ZipFile('/content/dogs-vs-cats.zip', 'r')
zip_ref.extractall('/content')
zip_ref.close()
```

In [7]:

```
import tensorflow as tf
from tensorflow import keras
from keras.layers import Dense, Conv2D, MaxPooling2D, Flatten, BatchNormalization, Dropout
from keras import Sequential
```

In [8]:

```
train_data = keras.utils.image_dataset_from_directory(
    directory = '/content/train',
    labels="inferred",
    label_mode="int",
    batch_size=32,
    image_size=(256, 256)
)

test_data = keras.utils.image_dataset_from_directory(
    directory = '/content/test',
    labels="inferred",
    label_mode="int",
    batch_size=32,
    image_size=(256, 256)
)
```

Found 20000 files belonging to 2 classes.

Found 5000 files belonging to 2 classes.

In [9]:

```
type(train_data)
```

Out[9]:

tensorflow.python.data.ops.prefetch_op._PrefetchDataset

In [10]:

```
for images, labels in train_data.take(1):
```

```
print("Batch of images:", images.shape)
print("Batch of labels:", labels)
```

```
Batch of images: (32, 256, 256, 3)
Batch of labels: tf.Tensor([0 1 1 1 0 0 0 0 1 0 1 1 1 1 1 1 1 0 0 1 0 1 0 1 1 1 1 0 1 1
0 1], shape=(32,), dtype=int32)
```

In [11]:

```
# we need to normalize our image data
def normalize(image,label):
    image = tf.cast(image/255, tf.float32)
    return image, label
```

In [12]:

```
train_data = train_data.map(normalize)
```

In [13]:

```
test_data = test_data.map(normalize)
```

In [14]:

```
model = Sequential()

model.add(Conv2D(32, kernel_size=(3,3),padding='valid', activation='relu', input_shape=(
256,256,3)))
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))

model.add(Conv2D(64, kernel_size=(3,3), padding='valid', activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))

model.add(Conv2D(128, kernel_size=(3,3),padding='valid', activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))

model.add(Flatten())

model.add(Dense(128, activation='relu'))
model.add(Dense(64, activation='relu'))
model.add(Dense(1,activation='sigmoid'))
```

In [15]:

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 254, 254, 32)	896
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_1 (Conv2D)	(None, 125, 125, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_2 (Conv2D)	(None, 60, 60, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 128)	0
flatten (Flatten)	(None, 115200)	0
dense (Dense)	(None, 128)	14745728
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 1)	1

```

=====
Total params: 14847297 (56.64 MB)
Trainable params: 14847297 (56.64 MB)
Non-trainable params: 0 (0.00 Byte)
=====

```

In [16]:

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

In [17]:

```
history = model.fit(train_data, epochs=8, validation_data=test_data)
```

```

Epoch 1/8
625/625 [=====] - 89s 132ms/step - loss: 0.6248 - accuracy: 0.63
90 - val_loss: 0.5241 - val_accuracy: 0.7404
Epoch 2/8
625/625 [=====] - 78s 124ms/step - loss: 0.4811 - accuracy: 0.76
49 - val_loss: 0.4617 - val_accuracy: 0.7838
Epoch 3/8
625/625 [=====] - 70s 110ms/step - loss: 0.3722 - accuracy: 0.83
25 - val_loss: 0.5073 - val_accuracy: 0.7684
Epoch 4/8
625/625 [=====] - 53s 84ms/step - loss: 0.2393 - accuracy: 0.898
4 - val_loss: 0.5847 - val_accuracy: 0.7730
Epoch 5/8
625/625 [=====] - 57s 91ms/step - loss: 0.1258 - accuracy: 0.951
3 - val_loss: 0.8162 - val_accuracy: 0.7818
Epoch 6/8
625/625 [=====] - 55s 87ms/step - loss: 0.0744 - accuracy: 0.973
0 - val_loss: 1.2892 - val_accuracy: 0.7776
Epoch 7/8
625/625 [=====] - 57s 90ms/step - loss: 0.0491 - accuracy: 0.985
1 - val_loss: 1.2086 - val_accuracy: 0.7816
Epoch 8/8
625/625 [=====] - 57s 90ms/step - loss: 0.0475 - accuracy: 0.984
1 - val_loss: 1.2306 - val_accuracy: 0.7880

```

In [18]:

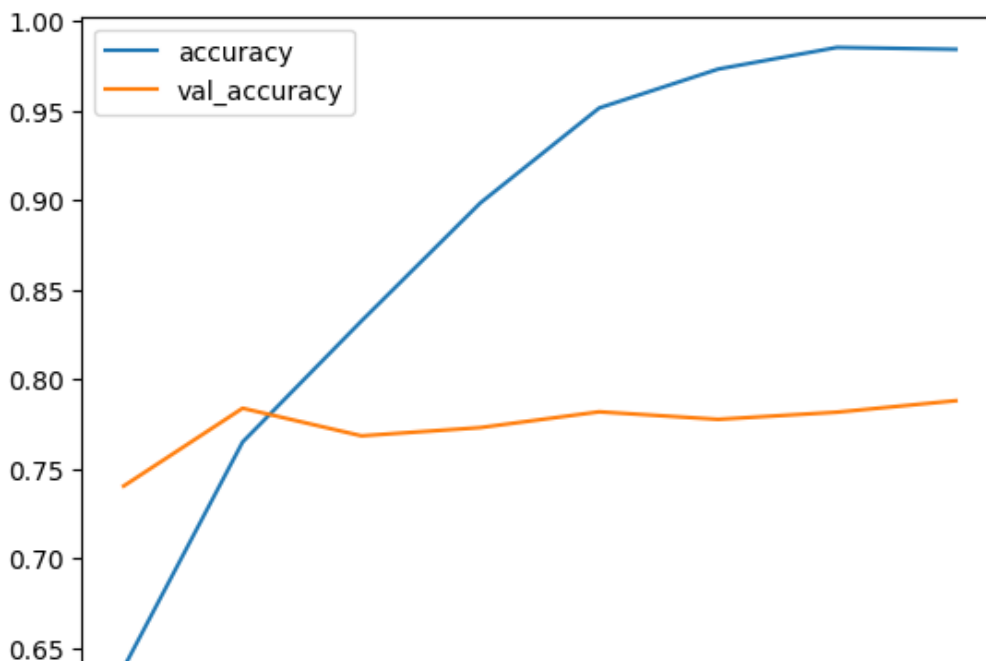
```


plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.legend()

```

Out[18]:

<matplotlib.legend.Legend at 0x7dd58fc1bb80>






In [19]:

```
plt.plot(history.history['loss'], label='loss')
plt.plot(history.history['val_loss'], label='val_loss')
plt.legend()
```

Out[19]:

<matplotlib.legend.Legend at 0x7dd53c107d60>



In [20]:

```
# overfitting
# 1--> Try to reduce overfitting using Dropout layer
```

In [21]:

```
# model1 = Sequential()

# model1.add(Conv2D(32, kernel_size=(3,3),padding='valid', activation='relu', input_shape
#=(256,256,3)))
# model1.add(Dropout(0.3))
# model1.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))

# model1.add(Conv2D(64, kernel_size=(3,3), padding='valid', activation='relu'))
# model1.add(Dropout(0.25))
# model1.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))

# model1.add(Conv2D(128, kernel_size=(3,3),padding='valid', activation='relu'))
# model1.add(Dropout(0.2))
# model1.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))

# model1.add(Flatten())

# model1.add(Dense(128, activation='relu'))
# model1.add(Dense(64, activation='relu'))
# model1.add(Dense(1,activation='sigmoid'))
```

In [22]:

```
# model1.compile(
#     optimizer='adam',
```

```
# loss='binary_crossentropy',
# metrics=['accuracy']
# )
```

In [23]:

```
# history1 = model1.fit(train_data,epochs=10, validation_data=test_data)
```

In [24]:

```
import cv2
```

In [25]:


```
test_img1 = cv2.imread('/content/cat.jpg')
```

In [26]:

```
plt.imshow(test_img1)
```

Out[26]:

<matplotlib.image.AxesImage at 0x7dd4b41dac80>



In [27]:

```
image1 = cv2.imread('/content/cat1.jpeg') #cat image
image2 = cv2.imread('/content/dog2.jpeg') # dog image
```

In [28]:

```
image1 = cv2.resize(image1, (256,256))
image2 = cv2.resize(image2, (256,256))
```

In [29]:

```
image1 = image1.reshape((1,256,256,3))
image2 = image2.reshape((1,256,256,3))
```

In [31]:

```
prediction_cat = model.predict(image1)
prediction_dog = model.predict(image2)

print("Prediction of cat image is :", prediction_cat[0][0])
```

```
print("Prediction of dog image is :", prediction_dog[0][0])
```

```
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 45ms/step
Prediction of cat image is : 0.0
Prediction of dog image is : 1.0
```

In [32]:

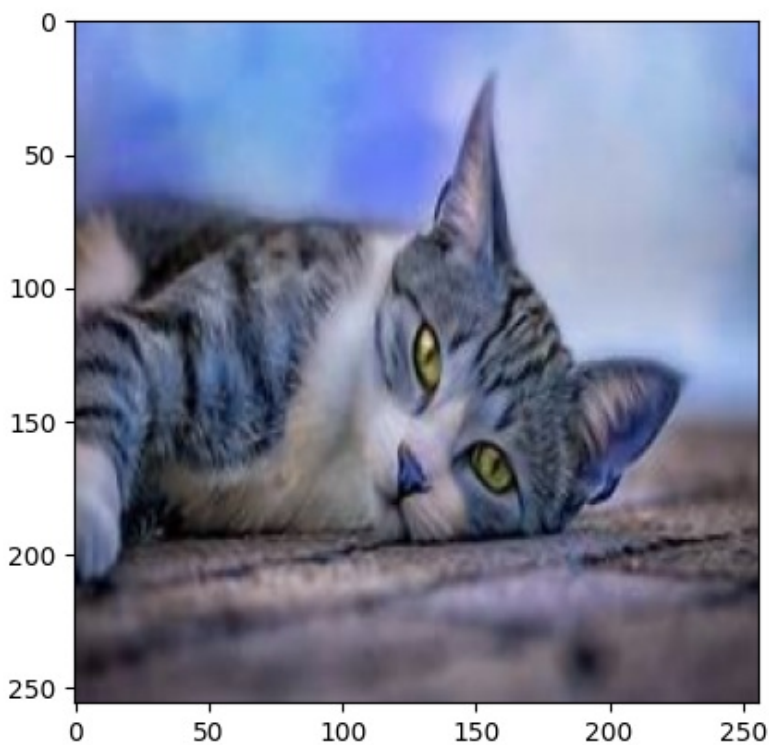
```
def Cat_or_Dog(image):
    image = cv2.resize(image, (256,256))
    temp_image = image
    image = image.reshape((1,256,256,3))
    prediction = model.predict(image)[0][0]
    plt.imshow(temp_image)

    if prediction == 1:
        print("The Image in the picture is DOG")
    else:
        print("The Image in the picture is CAT🐱")
```

In [33]:

```
image = cv2.imread('/content/cat2.jpeg')
Cat_or_Dog(image)
```


```
1/1 [=====] - 0s 20ms/step
The Image in the picture is CAT🐱
```




In [34]:

```
image = cv2.imread('/content/cat3.jpeg')
Cat_or_Dog(image)
```

```
1/1 [=====] - 0s 18ms/step
The Image in the picture is CAT🐱
```






In [35]:

```
image = cv2.imread('/content/dog3.jpg')  
Cat_or_Dog(image)
```

1/1 [=====] - 0s 34ms/step
The Image in the picture is DOG




In [37]:

```
image = cv2.imread('/content/dim.jpeg')  
Cat_or_Dog(image)
```

1/1 [=====] - 0s 22ms/step
The Image in the picture is DOG





In [38]:

```
import pickle  
pickle.dump(model, open('model.pkl', 'wb'))
```

In []: