

In [1]:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

In [2]:

```
!mkdir -p ~/.kaggle
!mv kaggle.json ~/.kaggle/
```

In [3]:

```
!kaggle datasets download -d salader/dogs-vs-cats
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /root/.kaggle/kaggle.json'
Downloading dogs-vs-cats.zip to /content
100% 1.06G/1.06G [00:46<00:00, 25.7MB/s]
100% 1.06G/1.06G [00:46<00:00, 24.6MB/s]

In [4]:

```
import zipfile
zip_ref = zipfile.ZipFile('/content/dogs-vs-cats.zip', 'r')
zip_ref.extractall('/content')
zip_ref.close()
```

In [5]:

```
import tensorflow as tf
from tensorflow import keras
from keras.layers import Dense, Conv2D, MaxPooling2D, Flatten, BatchNormalization, Dropout
from keras import Sequential
```

In [6]:

```
train_data = keras.utils.image_dataset_from_directory(
    directory = '/content/train',
    labels="inferred",
    label_mode="int",
    batch_size=32,
    image_size=(256, 256)
)

test_data = keras.utils.image_dataset_from_directory(
    directory = '/content/test',
    labels="inferred",
    label_mode="int",
    batch_size=32,
    image_size=(256, 256)
)
```

Found 20000 files belonging to 2 classes.
Found 5000 files belonging to 2 classes.

In [7]:

```
type(train_data)
```

Out[7]:

```
tensorflow.python.data.ops.prefetch_op._PrefetchDataset
```

In [8]:

```
for images, labels in train_data.take(1):  
    print("Batch of images:", images.shape)  
    print("Batch of labels:", labels)
```

Batch of images: (32, 256, 256, 3)

Batch of labels: tf.Tensor([0 0 0 0 0 0 1 0 1 0 1 1 0 0 0 1 1 0 0 1 1 0 1 0 1 1 1
0 1 1 1 1], shape=(32,), dtype=int32)

In [9]:

```
# we need to normalize our image data  
def normalize(image, label):  
    image = tf.cast(image/255, tf.float32)  
    return image, label
```

In [10]:

```
train_data = train_data.map(normalize)
```

In [11]:

```
test_data = test_data.map(normalize)
```

In [12]:

```
model = Sequential()  
  
model.add(Conv2D(32, kernel_size=(3,3),padding='valid', activation='relu',  
input_shape=(256,256,3)))  
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))  
  
model.add(Conv2D(64, kernel_size=(3,3), padding='valid', activation='relu'))  
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))  
  
model.add(Conv2D(128, kernel_size=(3,3),padding='valid', activation='relu'))  
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))  
  
model.add(Flatten())  
  
model.add(Dense(128, activation='relu'))  
model.add(Dense(64, activation='relu'))  
model.add(Dense(1,activation='sigmoid'))
```

In [13]:

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 254, 254, 32)	896
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_1 (Conv2D)	(None, 125, 125, 64)	18496

max_pooling2d_1 (MaxPoolin g2D)	(None, 62, 62, 64)	0
conv2d_2 (Conv2D)	(None, 60, 60, 128)	73856
max_pooling2d_2 (MaxPoolin g2D)	(None, 30, 30, 128)	0
flatten (Flatten)	(None, 115200)	0
dense (Dense)	(None, 128)	14745728
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 1)	65

=====

Total params: 14847297 (56.64 MB)
 Trainable params: 14847297 (56.64 MB)
 Non-trainable params: 0 (0.00 Byte)

In [14]:

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

In [15]:

```
history = model.fit(train_data, epochs=10, validation_data=test_data)
```


```
Epoch 1/10
625/625 [=====] - 69s 98ms/step - loss: 0.6009 - accuracy
: 0.6700 - val_loss: 0.5418 - val_accuracy: 0.7196
Epoch 2/10
625/625 [=====] - 62s 99ms/step - loss: 0.4693 - accuracy
: 0.7735 - val_loss: 0.4923 - val_accuracy: 0.7664
Epoch 3/10
625/625 [=====] - 54s 85ms/step - loss: 0.3603 - accuracy
: 0.8379 - val_loss: 0.4809 - val_accuracy: 0.8004
Epoch 4/10
625/625 [=====] - 54s 86ms/step - loss: 0.2259 - accuracy
: 0.9068 - val_loss: 0.6294 - val_accuracy: 0.8002
Epoch 5/10
625/625 [=====] - 56s 89ms/step - loss: 0.1207 - accuracy
: 0.9525 - val_loss: 0.8532 - val_accuracy: 0.7906
Epoch 6/10
625/625 [=====] - 57s 91ms/step - loss: 0.0790 - accuracy
: 0.9718 - val_loss: 1.0341 - val_accuracy: 0.7766
Epoch 7/10
625/625 [=====] - 53s 85ms/step - loss: 0.0631 - accuracy
: 0.9785 - val_loss: 1.0482 - val_accuracy: 0.7802
Epoch 8/10
625/625 [=====] - 56s 88ms/step - loss: 0.0455 - accuracy
: 0.9844 - val_loss: 1.1185 - val_accuracy: 0.7732
Epoch 9/10
625/625 [=====] - 54s 86ms/step - loss: 0.0332 - accuracy
: 0.9878 - val_loss: 1.4287 - val_accuracy: 0.7884
Epoch 10/10
625/625 [=====] - 58s 92ms/step - loss: 0.0309 - accuracy
: 0.9903 - val_loss: 1.5160 - val_accuracy: 0.7866
```

In [16]:

```
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.legend()
```

Out[16]:

<matplotlib.legend.Legend at 0x7ae5dd58f790>

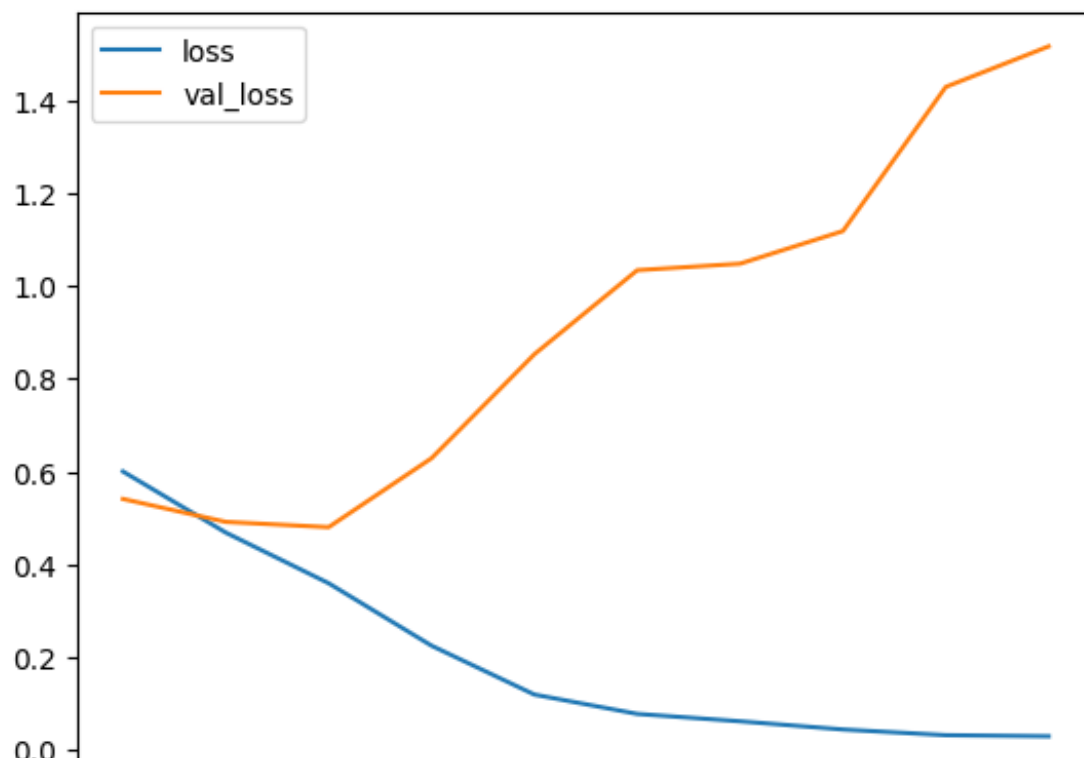


In [17]:

```
plt.plot(history.history['loss'], label='loss')
plt.plot(history.history['val_loss'], label='val_loss')
plt.legend()
```

Out[17]:

<matplotlib.legend.Legend at 0x7ae5dc272050>



In []:

```
# overfitting
# 1-->> Try to reduce overfitting using Dropout layer
```

In []:

```
# model1 = Sequential()

# model1.add(Conv2D(32, kernel_size=(3,3),padding='valid', activation='relu', input_shape=(256,256,3)))
# model1.add(Dropout(0.3))
# model1.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))

# model1.add(Conv2D(64, kernel_size=(3,3), padding='valid', activation='relu'))
# model1.add(Dropout(0.25))
# model1.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))

# model1.add(Conv2D(128, kernel_size=(3,3),padding='valid', activation='relu'))
# model1.add(Dropout(0.2))
# model1.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))

# model1.add(Flatten())

# model1.add(Dense(128, activation='relu'))
# model1.add(Dense(64, activation='relu'))
# model1.add(Dense(1,activation='sigmoid'))
```

In []:

```
# model1.compile(
#     optimizer='adam',
#     loss='binary_crossentropy',
#     metrics=['accuracy']
# )
```

In []:

```
# history1 = model1.fit(train_data,epochs=10, validation_data=test_data)
```

In [18]:

```
import cv2
```

In [19]:


```
test_img1 = cv2.imread('/content/cat1.jpeg')
test_img2 = cv2.imread('/content/cat2.jpeg')
```


In [20]:

```
plt.imshow(test_img1)
```

Out[20]:

<matplotlib.image.AxesImage at 0x7ae54813ba00>





In [21]:

```
image = cv2.imread('/content/dog1.jpeg')
```

In [22]:

```
image = cv2.resize(image, (256,256))
```

In [23]:

```
image = image.reshape((1,256,256,3))
```

In [24]:

```
prediction = model.predict(image)[0][0]  
prediction
```

1/1 [=====] - 0s 476ms/step

Out[24]:

1.0

In [31]:


```
def Cat_or_Dog(image):  
    image = cv2.resize(image, (256,256))  
    temp_image = image  
    image = image.reshape((1,256,256,3))  
    prediction = model.predict(image)[0][0]  
    plt.imshow(temp_image)  
  
    if prediction == 1:  
        print("The Image in the picture is DOG")  
    else:  
        print("The Image in the picture is CAT🐱")
```

In [33]:

```
image = cv2.imread('/content/cat2.jpeg')  
Cat_or_Dog(image)
```

1/1 [=====] - 0s 18ms/step

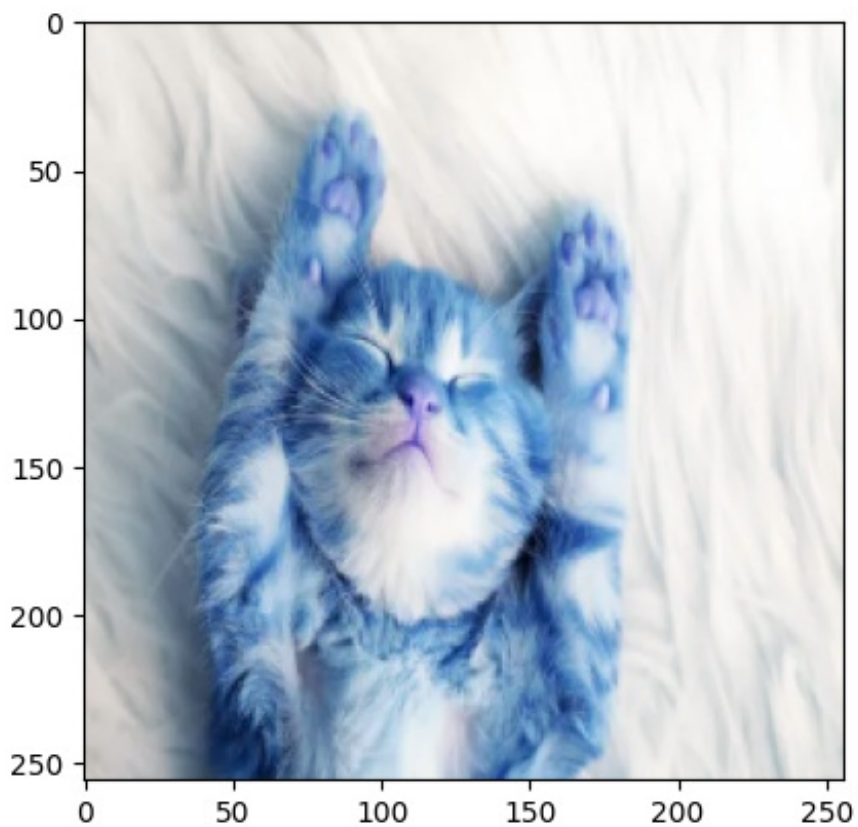
The Image in the picture is CAT🐱



In [34]:

```
image = cv2.imread('/content/cat3.jpeg')  
Cat_or_Dog(image)
```

1/1 [=====] - 0s 17ms/step
The Image in the picture is CAT🐱



In [35]:

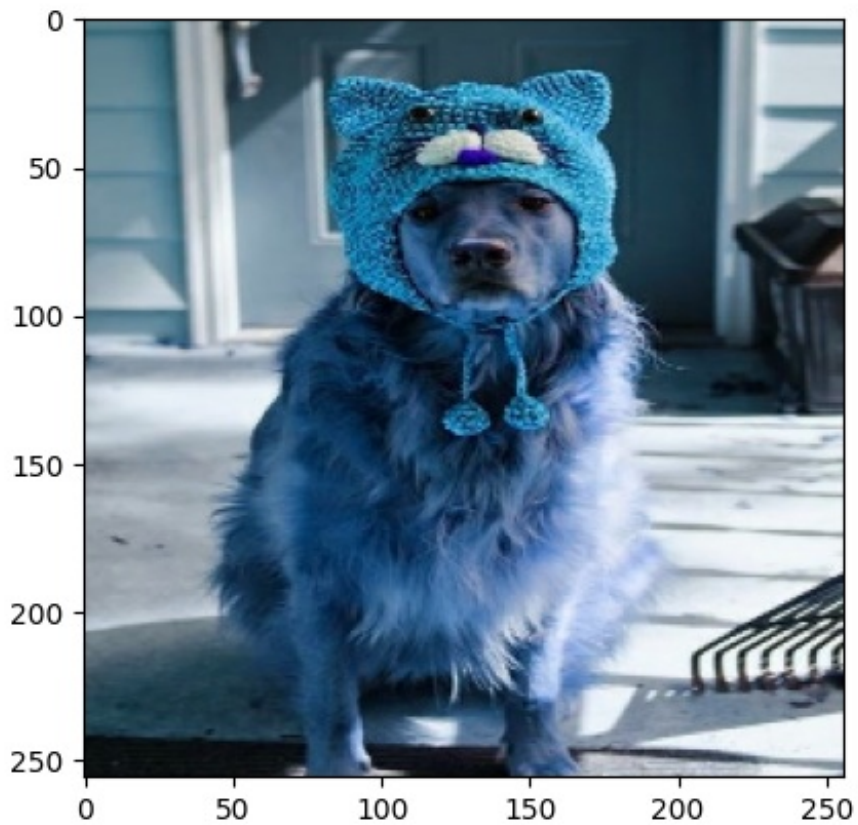
```
image = cv2.imread('/content/dog3.jpg')
```



```
Cat_or_Dog(image)
```

```
1/1 [=====] - 0s 18ms/step
```

```
The Image in the picture is DOG☐☐
```

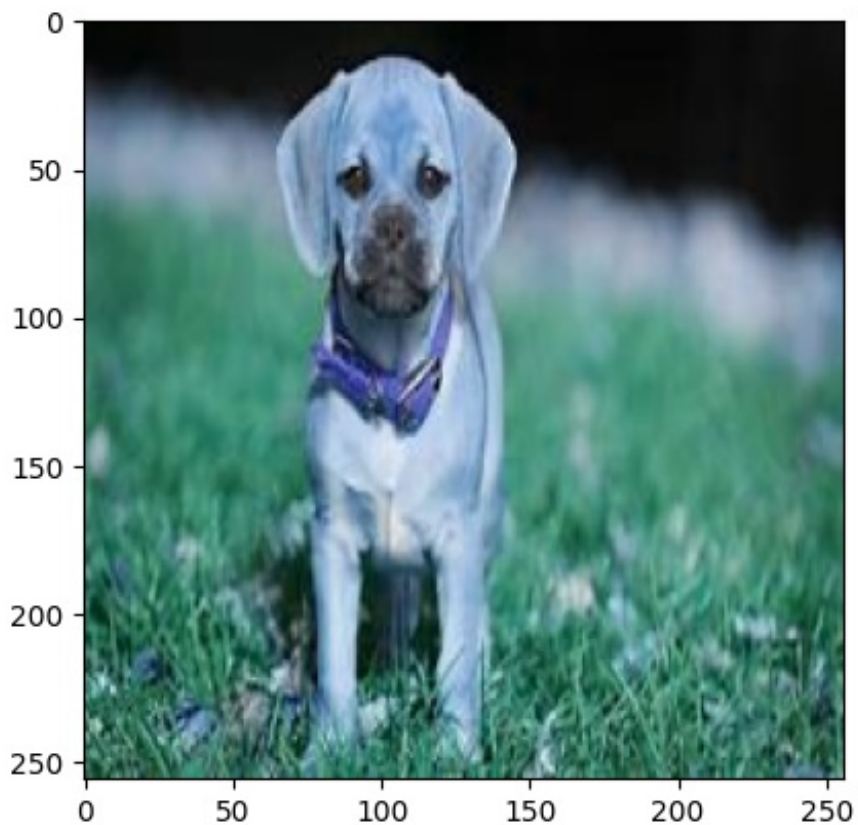


```
In [36]:
```

```
image = cv2.imread('/content/dog2.jpeg')  
Cat_or_Dog(image)
```

```
1/1 [=====] - 0s 18ms/step
```

```
The Image in the picture is CAT☐🐱
```



In [37]:

```
import pickle  
pickle.dump(model, open('model.pkl', 'wb'))
```

In []: